# Using the AWS command line interface to launch an EC2 server

## Ronald (Ryy) Glenn Thomas 2025-05-13

## Table of contents

1	Introduction	1
2	Scripts  2.1 Create security group script	4
3	Appendix 1 Set up AWS IAM	7
4	Appendix. Sample work session	8
5	Appendix. Undo	9
	5.1 Prerequisites	
	5.2 Step-by-Step Implementation	
	5.3 Key Takeaways	9
	5.4 Further Reading	9

## 1 Introduction

This post suggests a straightforward strategy for quickly getting a secure AWS EC2 server up and running.

In a separate post (here) we address the same task using the interactive Elastic Compute Cloud (EC2) dashboard. While its certainly instructive to use the EC2 console interface to set up a working environment in AWS and launch a custom server, it can become a tedious process after the first few repetitions. In this post we'll present a set of bash shell scripts to perform the same task, making use of the AWS command line interface (CLI).

To get started, install and configure the awscli app using the commands (on your mac workstation):



```
> brew install awscli
> aws configure
```

The aws app will initially ask for your IAM credentials. If you don't have an IAM ID, Appendix 1 (here) provides details on obtaining IAM credentials from the AWS interface. Additional instructions from Amazon for installing the AWS command line interface can be found (here).

Instructions to install the homebrew software management system on a mac can be found (here).

## 2 Scripts

Nine parameters (besides the IAM ID) are required to be set for automated instance (server) generation via the aws API. The first eight are likely to be static and we suggest you "hardcode" them as environment variables in your shell configuration file.

An example follows (with z-shell syntax):

Parameters 1) and 2) vpc\_id and subnet\_id refer to TBD. The default vpc\_id is assigned by AWS and can be found on the EC2 dashboard. The default VPC will have two subnets. Select one from the EC2 dashboard in which to launch the instance (subnet\_id).

Add the following lines to your .zshrc file:

```
> export vpc_id="vpc-14814b73"
> export subnet_id="subnet-f02c90ab"
```

Parameters 3), 4) and 5) ami\_id, storage\_size and instance\_type define the operating system and the capabilities of the server.

Add the following three lines to your .zshrc file:

```
export ami_id="ami-014d05e6b24240371"
export instance_type="t2.micro"
export storage_size="30"
```

Parameters 6) and 7) key\_name and security\_grp identify the ssh key-pair and the firewall:

```
export key_name="power1_app"
export security_grp="sg-0fef542d93849669c"
```

Parameter 8) is the static\_ip that identifies the server on the web:

```
export static_ip="13.57.139.31"
```

A ninth parameter proj\_name can also be hardcoded or supplied at the time the script is called. Below we offer four bash scripts.

1) The first creates a key pair to allow encrypted ssh communication between the server and your workstation.

- 2) The second generates a security group for the virtual server, i.e. a firewall.
- 3) The third script generates the virtual server taking physical server location, instance characteristics, firewall, static IP and project name as parameters.
- 4) The fourth script installs required software following server launch.

## 2.1 Create security group script

Generate security group: '

Example:

```
> aws_create_security_group.sh -s $proj_name -g -k
```

This version of the script creates a security group with options to open ports: 22, 80, 3838, 443, 9000, and 9001 with flags -g, -i, -j, -k, -l, -m respectively.

The default, i.e. no flags set, is to open 22 and 443 only. the security group name is set as the base directory name.

```
#!/usr/bin/env bash
Help()
{
echo The script generates a new security group
echo the group name is given with the -n flag.
echo ports are specificed with the -p flag. Any number of ports can be listed
echo Anticipated incoming ports are 22 ssh, 80 http, 3838 shiny and 443 https.
echo Script will fail if group name is already in use on EC2.
echo reads vpc_id from the environment variables set in .zshrc
echo example usage for ports 22, 80 and 443:
echo aws_create_security_group.sh -n power1_app -p 22 -p 80 -p 443
sg_grp_name=`basename $PWD`
while getopts ":hp:n:" opt; do
    case $opt in
        p ) ports+=("$OPTARG") ;; # use the split+glob operator
        n ) sg_grp_name=$OPTARG ;;
       h ) Help
            exit ;;
        * ) echo 'error in command line parsing. Expect options n and p' >&2
    esac
done
echo "sg group name = $sg_grp_name"
```

```
aws ec2 create-security-group \
   --group-name $sg_grp_name \
    --description "security group" \
    --tag-specifications \
   "ResourceType=security-group, Tags=[{Key=Name, Value=$sg_grp_name}]" \
    --vpc-id $vpc_id > temp.txt
wait
security_grp=`jq -r .GroupId temp.txt`
echo "security group ID = $security_grp"
for i in "${ports[@]}"
  aws ec2 authorize-security-group-ingress \
   --group-id $security_grp \
   --protocol tcp \
   --port ${i} \
    --cidr "0.0.0.0/0" > /dev/null
 done
```

## 2.2 Create new key pair with a project name flag

Example usage: Note run with one parameter for optional keypair name.

```
> aws_create_keypair.sh -k $keypair_name
```

```
#!/usr/bin/env bash
Help()
{
echo The script generates a new keypair
echo the keypair name is given with the -k flag.
echo Script will fail if pair name is already in use on EC2.
echo aws_create_keypair.sh -k power1_app
}
while getopts 'hk:' flag; do
  case "${flag}" in
   h) Help
      exit;;
   k) key_pair_name=${OPTARG};;
  esac
done
base=`basename $PWD`
if [ -z "$key_pair_name" ]
then
 key_pair_name=$base
```

```
fi
echo "key_pair_name is $key_pair_name"

cd ~/.ssh
rm -f ~/.ssh/$key_pair_name.pem
aws ec2 create-key-pair --key-name $key_pair_name \
    --query 'KeyMaterial' --output text > ~/.ssh/$key_pair_name.pem

wait
chmod 400 ~/.ssh/$key_pair_name.pem
```

#### 2.3 Generate instance

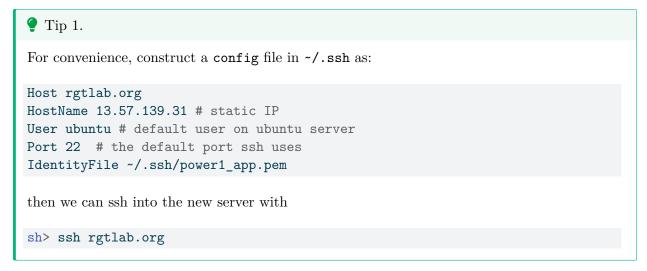
start up script. > aws\_create\_instance.sh -p power1\_app

```
#!/usr/bin/env bash
Help()
{
echo "Notes on currect parameters:"
echo "security group should be in place already. check on EC2. If not,
run ./awscli_create_security.sh. "
echo "Key pair should be in place. check on EC2 and in ~/.ssh.
If not run ./create_keypair.sh. "
echo "ami id is for ubuntu linux 22.04 LTS.
If not what is desired check EC2 list of instances."
echo "Check static IP: nslookup IPaddress.
Should point to the domain name e.g. rgtlab.org "
echo Usage: >aws_create_instance.sh -p power1_app
echo ""
echo "Review parameters: "
echo "---"
echo "proj_name is $proj_name"
echo "keypair_name is $keypair_name"
echo "vpc_id: $vpc_id";
echo "subnet_id: $subnet_id";
echo "ami_id: $ami_id";
echo "security_grp: $security_grp";
echo "static_ip: $static_ip";
echo "type: $type";
echo "size: $size";
}
while getopts 'hp:' flag; do
  case "${flag}" in
   h) Help
      exit;;
```

```
p) proj_name=${OPTARG};;
  esac
done
base=`basename $PWD`
if [ -z "$proj_name" ]
then
 proj_name=$base
fi
aws ec2 run-instances \
--image-id $ami_id \
--count 1 \
--instance-type $instance_type \
--key-name $keypair_name \
--security-group-ids $security_grp \
--subnet-id $subnet_id \
--block-device-mappings "[{\"DeviceName\":\"/dev/sda1\",\"Ebs\":{\"VolumeSize\":$storage_size}
--tag-specifications "ResourceType=instance, Tags=[{Key=Name, Value=$proj_name}]" \
--user-data file://~/Dropbox/prj/c060/aws_startup_code.sh
iid0=`aws ec2 describe-instances --filters "Name=tag:Name, Values=$proj_name" | \
    jq -r '.Reservations[].Instances[].InstanceId'`
echo $iid0
read -p "enter instance id:" iid
echo "instance id: $iid"
aws ec2 associate-address --public-ip $static_ip --instance-id $iid
```

#### aws\_startup.sh

```
#!/bin/bash
apt update
# Add Docker and Docker Compose support to the Ubuntu's packages list
apt-get install curl -y
apt-get install gnupg -y
apt-get install ca-certificates -y
apt-get install lsb-release -y
sudo install -m 0755 -d /etc/apt/keyrings
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | \
  sudo gpg --dearmor -o /etc/apt/keyrings/docker.gpg
sudo chmod a+r /etc/apt/keyrings/docker.gpg
echo "deb [arch="$(dpkg --print-architecture)" \
signed-by=/etc/apt/keyrings/docker.gpg] https://download.docker.com/linux/ubuntu \
"$(. /etc/os-release && echo "$VERSION_CODENAME")" stable" | \
sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
apt-get update
apt-get install docker-ce docker-ce-cli containerd.io docker-compose-plugin -y
su ubuntu -
usermod -aG docker ubuntu
```



Change the access permissions: sudo chmod 600 power1ssh.pem to be more restrictive.

## 3 Appendix 1 Set up AWS IAM

This appendix provides details on how to initiate batch processing via the AWS CLI desktop application.

From aws web site:

Log into the AWS console.

Search for IAM service. Navigate to IAM dashboard.

Select User groups. Create a user group based on the Power User profile.

Call it admin. Include ryy\_iam in the admin group.

Select Users in left hand panel.

Then select Create User button (in upper right).

Then enter a User name in the form, say zenn. Click Next (lower right)

Then Create User.

Click on the user name

In the page that comes up. Select Security Credentials tab (center of page).

Under Access Keys panel click Create access key (left side or bottom of panel).

Click Command Line Interface CLI

and at the bottom of the page click the checkbox "I understand...".

Finally select Create access key and

choose Download .csv file (lower right).

Navigate Download screen to local ~/.aws directory.(may need shift-cmd-. on mac)



#### Click Done

Now in the terminal on your workstation, configure the aws cli app via the command.

```
> aws configure
```

Using cut and paste enter info from the credentials file just downloaded. After entering the AWS Access Key ID and AWS Secret Access Key information you will be asked for a Region, (our region is us-west-1), and an output format (suggested output format is JSON).

"AWS
Identity and
Access
Management
(IAM) is a
web service
that helps
you securely
control access
to AWS
resources."

Tip - In the

AWS Console, in the

navigation

pane, select your VPCs.

The VPC page

lists the vpc\_id and the subnets

subnet\_id.

## 4 Appendix. Sample work session

Start from scratch. Assume the following:

- 1. aws cli is configured.
- 2. no security group has been defined
- 3. no key pair has been generated
- 4. vpc ID and subnet ID are known and stored as environment variables.
- 5. project name is power1\_app

In the following we'll spin up a ubuntu server (AMI) with type t2-micro (1 vCPU and 1gb memory) and 30 GB (size) hard drive.

step 1. Generate a security group named power1\_app, and get the security\_grp from the script output and store it as an environment variable.

```
> aws_create_security_group.sh -n power1_app -p 22 -p 80 -p 443
```

step 2. add key pair with name power1\_app

```
aws_create_keypair.sh -k power1_app
```

step 3. get a new elastic IP address and add it to z-shell configuration file. and modify ssh config file to add the IP address and the ssh private key name. if new IP is: 204.236.167.50

```
sed -i '.bak' '/static/d' ~/.config/zsh/.zsh_export
sed -i '.bak' '/security/d' ~/.config/zsh/.zsh_export
echo "export static_ip='204.236.167.50'" >> ~/.config/zsh/.zsh_export
echo "export security_grp='sg-0fda72c2879d6b2ad'" >> ~/.config/zsh/.zsh_export
```

Generate instance:

```
aws_create_instance.sh -p power1_app
```

```
cd ~/.ssh
sed -i '.bak' '/HostName/d' config
echo "export HostName='204.236.167.50'" >> config
```

## 5 Appendix. Undo

- To remove the AWS instance and Gitlab elements of a project do the following:
- log into AWS/EC2 console
- Terminate instance (Instance state menu)
- delete security group (Network & Security tab, Actions menu)
- release IP address (Network & Security tab, Actions menu)
- delete SSH key pair (Network & Security tab, Actions menu)
- Log into Gitlab (gitlab.com)
- Delete project(s) (project/settings/General/advanced)

## 5.1 Prerequisites

In development

## 5.2 Step-by-Step Implementation

In development

## 5.3 Key Takeaways

In development

## 5.4 Further Reading

In development