

Set up a virtual server on AWS (in anticipation of hosting Shiny apps)

Ronald (Ryy) Glenn Thomas

Table of contents

Introduction	1
Hosting	2
Select a hosting service	2
AWS Working Environment	3
Ssh key pair	4
Firewall	5
Static IP address	5
Domain Name	5
Select and launch instance	6
Access server	6
Appendix: Tips	7
Appendix: Non-opinionated instructions	7

Introduction

Ok! I've got my shiny app running just the way I want it.
Works great! Now how do I get it up on the web and shared
with my client?



Photo by Nathan Waters on
Unsplash

Assuming we have a working shiny app, we often next need to address the task of hosting the app on a (virtual) server to share with our collaborators. There are many ways to accomplish this. In this post, we'll describe how to 'spin up' a server on Amazon Web Service EC2. In a future post we'll show how, in just a few steps, through the application of Docker, R, Shiny, and Caddy (webserver) functionality we can have a fully functional and secure app available on the web to share with colleagues.

Hosting

In order to host a shiny app, say `power1_shiny`, online we'll need to complete the following tasks:

1. create a virtual server (connected via ssh) with a firewall
2. obtain a static IP address (to identify the server online)
3. obtain a domain name (name for IP address)
4. install and configure a webserver (tool to interact with https protocol requests and respond)
5. obtain and install an SSL certificate (to allow encrypted communication)
6. setup an authentication method (password protection)
7. configure a reverse proxy method (translate https (port 443) requests to Shiny (port 3838))

In this post we'll address the first three of these. Tasks four through seven will be accomplished using docker-compose and Caddy, as described in the next post.

At first glance these 7 requirements can appear daunting, but on closer inspection all can be met with relative ease and minimal cost (using a cloud-hosting service, e.g. Amazon's EC2 or Digital Ocean, and a "leased" domain name from, e.g. GoDaddy, or Amazon's Route 53) or at no cost if we have our own server with IP address, and domain name.

Select a hosting service

There are a number of cloud based server options to choose from: Microsoft Azure, Oracle, Google Cloud, Amazon AWS

EC2, Digital Ocean to name a few. Each has their own approach to setting up a custom virtual server. Several have free or low-cost service tiers available.

In this power we'll provide a step-by-step description of the process with AWS EC2.


AWS is a reasonable choice for setting up a small custom server. Its not the cheapest option but the system is very well documented and, in our experience, reliable.

To start open the EC2 console visit the URL:

```
https://aws.amazon.com/console
```

Next choose regional service. For me its “N. California”.

Next create an account or sign in and navigate to the EC2 dashboard. Its through the dashboard that we'll define the parameters for the type of server to launch and the mechanisms for communicating with it.



AWS Working Environment

That is, along with selecting a server we'll need to set up a working environment. We recommend setting up the working environment before launching the server, as it saves some back and forth. This working environment consists of four main components:

1. A secure shell (ssh) key-pair to allow we to remotely and securely login to the virtual server once its launched.
2. A firewall to restrict access to only secure connections. The firewall closes off all incoming traffic except through those ports specifically named.
3. A static IP address. This is required for maintaining the link between the domain name and the server when re-booting. (The default is for the instance/server to be assigned a new IP each time its rebooted.) and

4. A domain name, say rgtlab.org. A domain name is not required but will facilitate collaborator access by not needing to use the IP address directly.

Ssh key pair

The first time we create an AWS account we need to exchange an ssh key pair with AWS. The pairs consist of a **private** and a **public** key. The pair of keys will allow us to login securely to any server we launch on EC2. We can generate an ssh key pair in one of two ways in EC2. Either, locally, on our workstation and upload the public key to EC2, or have EC2 generate the key pair and download the private key.

For the first option we create a directory on our workstation to hold the keys and navigate to it. e.g. `~/.ssh`. Generate the keys with the command

```
ssh-keygen -m PEM
```

“PEM” defines the key format. More information on public key authentication can be found [here](#). In the interactive dialog name the key prefix something like `power1_app_ssh`. The dialog will ask for a passphrase. Enter one for an additional level of security, but its not required. The `ssh-keygen` program will generate two files: `power1_app_ssh` and `power1_app_ssh.pub`

To complete the process return to the EC2 dashboard select **Keys Pairs** in the left panel. The select **Create key pair** button in the upper right of the page.

A form appears and asks for a name. Enter something like `power1_app_ssh`. Select **RSA** for key pair type and `.pem` for key file format to match the `ssh-keygen` process we went through on our workstation.

asks for the location of the public key. Browse to the `~/.ssh` directory on wer workstation and import the public key `power1_app_ssh.pub`.

Give the pair a name, say `power1_app_ssh`, and the keys will be created and the private key `power1_app_ssh.pem` will be downloaded to our local machine to the `~/.ssh` directory. Change the access permissions for the private key with the following command:

```
sudo chmod 600 power1_app_ssh.pem
```

Firewall

To create a firewall click on **Network settings** in the left hand panel. Choose **Create security group** and select **Allow SSH traffic** and **Allow HTTPS traffic**. This will create a firewall that leaves open only ports 22 and 443, for `ssh` and `https` incoming traffic respectively. The default name for the firewall will be something like `launch-wizard-6`. name the security group `power1_firewall`.

Static IP address

The next step is to use the `elastic` IP service to get a static IP that can be assigned to the server. Navigate to **Network and Security** again and select **Allocate Elastic IP**. An IP will be assigned from the EC2 pool of available IPv4 IP addresses.

Domain Name

To obtain a dedicated domain name go to Amazon route 53 to select a domain name and associate it with our static IP.

Once a domain name is acquired, e.g. `rgtlab.org`, associate it with our static IP address. This can be done via the **Route 53** service. To associate domain name `rgtlab.org` with elastic IP do the following in Route 53:

- click on `hosted zones` in side panel
- click on `rgtlab.org` in center panel
- click on checkbox for `rgtlab.org` type=A line
- then click on edit record in right panel

- change IP address to the static IP (e.g. 13.57.139.31).

Select and launch instance

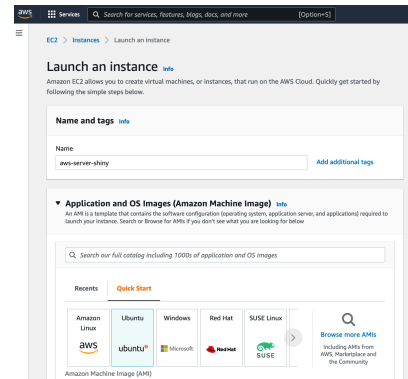
2. From “Quick Start” in the EC2 dashboard select an operating system for the server. Our recommendation is the Ubuntu OS. Ubuntu is based on linux. Click the Ubuntu button.
- Name the server, say power1
3. Next choose an instance **type**, e.g. “t2-micro”. Different instance types are combinations of, processors, memory, storage capacity, and network performance.
4. click **Configure Instance Details**
5. choose a Key pair (use power1_app_ssh.rsa from wer environment) .
6. Add security group, e.g. ‘power1_firewall’ from wer environment.
7. choose 30 GB of EBS General Purpose (SSD) or Magnetic storage. Thirty GBs is the maximum allowed in the ‘Free tier’ of servers on AWS. In our experience smaller disk sizes can lead to problems.
8. click **Launch Instance**

to Launch the Instance.

Access server

On wer laptop log into server with

```
ssh -i "~/ssh/power1_app_ssh.pem" ubuntu@rgtlab.org
```



Appendix: Tips

Tip 1.

For convenience, construct a config file in `~/.ssh` as:

```
Host rgtlab.org
HostName 13.57.139.31 # static IP
User ubuntu # default user on ubuntu server
Port 22 # the default port ssh uses
IdentityFile ~/.ssh/power1_app_ssh.rsa
```

then we can ssh into the new server with

```
sh> ssh rgtlab.org
```

Appendix: Non-opinionated instructions

Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>

From the top navigation bar, select a Region to create an instance in. For me its “N. California”.

Create an account or sign in and navigate to the EC2 dashboard.

In the left side panel select **Key Pairs** (under Network and Security).

At the top right select the **Create key pair** button. A **Key Pair** form will open.

Give the key pair a name. Something like `power1_app_ssh`. Select a key pair type, suggest RSA. Select a **Private key file format**, suggest `.pem`

Below the form select the **Create key pair** button. A pair of keys will be created and the private key `power1_app_ssh.pem`

will be downloaded to the local machine. In my case to the default `~/Downloads` directory.

Move the file to the `~/.ssh` directory: `mv ~/Downloads/power1_app_ssh.pem ~/.ssh`

Change the access permissions: `sudo chmod 600 power1ssh.pem` to be more restrictive.