

Research Compendia for Full Reproducibility in R: An rrtools, renv, and Docker Strategy

R.G. Thomas

2025-06-21

This white paper presents a comprehensive approach to achieving reproducibility in R workflows by combining three powerful tools: rrtools for creating structured research compendia, renv for R package management, and Docker for containerizing the computing environment. The rrtools package provides a standardized research compendium structure, renv manages package dependencies, and Docker ensures consistent execution environments. Together, these tools create self-contained research compendia that run identically across different systems. The paper includes a practical case study demonstrating multi-developer collaborative workflows with clear governance roles, where a project maintainer manages the technical infrastructure while multiple contributors extend the research analysis.

Table of contents

Executive Summary	3
Motivation	4
1 Introduction	5
1.1 The Challenge of Reproducibility in R	5
1.2 A Three-Level Solution	5
2 rrtools: Project-Level Reproducibility	6
2.1 What is rrtools?	6
2.2 Key Features of rrtools	6
2.3 Enhanced rrtools Workflow	7
2.4 Enhanced Research Compendium Structure	7
3 renv: Package-Level Reproducibility	8
3.1 What is renv?	8

3.2	Key Features of renv	8
3.3	Basic renv Workflow	8
4	Docker: System-Level Reproducibility	9
4.1	What is Docker?	9
4.2	Docker's Role in Reproducibility	9
4.3	Docker Components for R Workflows	9
5	Combining rrtools, renv, and Docker: A Comprehensive Approach	11
5.1	Why Use All Three?	11
5.2	Integration Strategy with Governance Model	12
6	Practical Example: Collaborative Research Compendium Development with Testing	14
6.1	Project Scenario	15
6.2	Step-by-Step Implementation	15
6.3	Developer 1: Project Setup and Initial Analysis	15
6.4	Developer 2: Extending the Analysis	24
7	Continuous Integration Extension	32
7.1	Key Benefits Demonstrated in This Example	33
8	Best Practices and Considerations	33
8.1	When to Use This Approach	33
8.2	Tips for Efficient Implementation	33
8.3	Testing Strategies for R Analyses	34
8.4	Potential Challenges	34
9	Conclusion	34
10	References	35
Appendix A:	GitHub Personal Access Token Setup	35
10.1	Step-by-Step Token Creation	35
10.2	Token Security Best Practices	36
10.3	Alternative: Using GitHub CLI	36
10.4	Troubleshooting Common Issues	36
11	Appendix: Comprehensive Test Suite for Palmer Penguins Analysis	37
11.1	Test File: tests/testthat/test-comprehensive-analysis.R	37
11.2	Running the Tests	42
11.3	Test Categories Explained	42
Appendix B:	Enhanced Directory Structure	43
11.1	Key Features Explained:	44

Appendix C: Docker Workflow Options	44
11.1 Option 1: Make Commands (Recommended)	45
11.2 Option 2: Docker Compose Services	45
11.3 Option 3: Direct Docker Commands	46
11.4 Volume Mounting Strategies	48
11.5 Choosing the Right Approach	48
Appendix D: GitHub Actions CI/CD Setup	48
11.1 Understanding GitHub Actions for Research	49
11.2 Step-by-Step Setup	49
11.2.1 Step 1: Create Workflow Directory	49
11.2.2 Step 2: Docker-based CI Workflow	49
11.2.3 Step 3: R Package Check Workflow	50
11.2.4 Step 4: Automated Paper Rendering	52
11.2.5 Step 5: Container Registry Integration	53
11.3 Workflow Explanations	55
11.3.1 Docker CI Workflow Features:	55
11.3.2 R Package Check Features:	55
11.3.3 Paper Rendering Features:	55
11.3.4 Container Publishing Features:	55
11.4 Authentication and Permissions	55
11.4.1 Built-in GITHUB_TOKEN:	55
11.4.2 Setting Repository Permissions:	56
11.4.3 Using Personal Access Tokens (Advanced):	56
11.5 Integration with Collaborative Workflow	56
11.5.1 Pull Request Integration:	56
11.5.2 Branch Protection Rules:	56
11.6 Monitoring and Troubleshooting	56
11.6.1 Viewing Workflow Results:	56
11.6.2 Common Issues and Solutions:	57
11.6.3 Performance Optimization:	57

Executive Summary

Reproducibility is key to conducting professional data analysis, yet in practice, achieving it consistently with R workflows can be quite challenging. R projects frequently break when transferred between computers due to mismatched R versions, package dependencies, or inconsistent project organization. This white paper describes a comprehensive approach to solving this problem by combining three powerful tools: **rrtools** for creating structured research compendia, **renv** for R package management, and **Docker** for containerizing the computing

environment. Together, these tools ensure that an R workflow runs identically across different computers by providing standardized project structure, identical R packages and versions, consistent R versions, and the same operating system libraries as the original setup.

Motivation

Imagine you’ve written code that you want to share with a colleague. At first glance, this may seem like a straightforward task—simply share the files electronically. However, ensuring that your colleague can run the code without errors, and obtain the same results is often much more challenging than anticipated.

When sharing R code, several potential problems can arise that can lead to code that won’t run or won’t match your results:

- Different versions of R installed on each machine
- Mismatched R package versions
- Missing or mismatched system dependencies (like pandoc or LaTeX)
- Missing supplemental files referenced by the program (bibliography files, LaTeX preambles, datasets, images)
- Different R startup configurations (.Rprofile or .Renv)
- Different Operating Systems (macOS, Windows, Linux, etc.)

A real-world scenario often unfolds like this:

1. You email your analysis files to your colleague, Joe
2. Joe attempts to run your analysis with the commands you provided
3. But R isn’t installed on Joe’s system
4. After installing R, Joe gets an error: “could not find function ‘render’ ” since he doesn’t have the `rmarkdown` package installed
5. Joe installs the `rmarkdown` package and runs the R command again
6. Now pandoc is missing
7. After installing pandoc, a required package, say `ggplot`, is missing
8. After installing `ggplot`, several external files are missing (e.g. bibliography, images)
9. And so on...

This cycle of troubleshooting can be time-consuming and frustrating. Even when the code eventually runs, there’s no guarantee that Joe will get the same results that you did.

To ensure true reproducibility, your colleague should have a computing environment as similar to yours as possible. Given the dynamic nature of open source software, not to mention hardware and operating system differences, this can be difficult to achieve through manual installation and configuration.

The approach outlined in this white paper offers a more robust solution. Rather than sending standalone text files, with modest additional effort, you can provide a complete, containerized, hardware and OS independent environment that includes everything needed to run your analysis. With this approach, your colleague can run a simple command like:

```
docker run \  
-v "$(pwd):/home/analyst/project" \  
-v "$(pwd)/analysis/figures:/home/analyst/output" \  
ghcr.io/username/penguins_analysis
```

(The details of this docker command are explained below.)

This creates an identical R environment on their desktop, ready for them to run or modify your code with confidence that it will work as intended.

1 Introduction

1.1 The Challenge of Reproducibility in R

R has become a standard tool for data science and statistical analysis across numerous scientific disciplines. However, as R projects grow in complexity, they often develop complex webs of dependencies that can make sharing and reproducing analyses difficult. Some common challenges include:

- Different R versions across machines
- Incompatible package versions
- Missing system-level dependencies
- Operating system differences (macOS vs. Windows vs. Linux)
- Conflicts with other installed packages
- R startup files (.Rprofile, .Renviron, .RData) that can affect code behavior

These challenges often manifest as the frustrating “it works on my machine” problem, where analysis code runs perfectly for the original author but fails when others attempt to use it. This undermines the scientific and collaborative potential of R-based analyses.

1.2 A Three-Level Solution

To address these challenges comprehensively, we need to tackle reproducibility at three distinct levels:

1. **Project-level reproducibility:** Ensuring consistent project structure and organization using research compendium standards

2. **Package-level reproducibility:** Ensuring exact package versions and dependencies are maintained
3. **System-level reproducibility:** Guaranteeing consistent R versions, operating system, and system libraries

The strategy presented in this white paper leverages **rrtools** for project-level structure, **renv** for package-level consistency, and **Docker** for system-level consistency. When combined, they provide a robust framework for end-to-end reproducible R workflows with proper research compendium organization.

2 rrttools: Project-Level Reproducibility

2.1 What is rrttools?

rrtools is an R package developed by Ben Marwick that provides instructions, templates, and functions for creating research compendia suitable for reproducible research. A research compendium is a standard and easily recognizable way of organizing the digital materials of a research project to enable others to inspect, reproduce, and extend the research.

2.2 Key Features of rrttools

rrtools creates a structured research compendium that follows established conventions:

- **Standardized directory structure:** Creates organized folders for data, analysis, papers, and figures following research compendium best practices
- **R package framework:** Uses R package structure to leverage existing tools for dependency management, documentation, and testing
- **Integrated documentation:** Automatically generates README files, citation information, and licensing documentation
- **Docker integration:** Provides functions to create Dockerfiles specifically designed for research compendia
- **Publication-ready structure:** Creates templates for academic papers and reports using R Markdown/Quarto

2.3 Enhanced rrttools Workflow

The enhanced rrttools workflow using the custom setup script involves:

```
# Run the rrttools setup script in your project directory
./setup_rrtools.sh
```

This automated script creates a comprehensive research compendium that includes:

- **Enhanced R package structure** with proper DESCRIPTION, NAMESPACE, and documentation
- **Comprehensive directory organization** with data, analysis, scripts, and documentation folders
- **Automated renv setup** with a curated list of commonly-used R packages
- **Docker integration** using rocker/r-ver with TinyTeX support
- **GitHub Actions workflows** for automated testing, checking, and paper rendering
- **Make-based build system** supporting both native R and Docker workflows
- **Symbolic links** for easy navigation between directories

The script automatically organizes existing files and creates a professional research compendium structure that follows best practices for reproducible research.

2.4 Enhanced Research Compendium Structure

The enhanced rrttools setup creates a comprehensive directory structure that follows research compendium best practices. The structure includes organized data folders, analysis directories, testing frameworks, and automated workflows.

Key organizational principles:

- **Data management:** Separate folders for raw, derived, and external data with proper documentation
- **Analysis workflow:** Dedicated spaces for papers, figures, tables, and working scripts
- **Package structure:** Professional R package organization with documentation and testing
- **Automation support:** Integration with Docker, GitHub Actions, and build systems

For the complete directory structure and detailed explanations, see Appendix B: Enhanced Directory Structure.

3 renv: Package-Level Reproducibility

3.1 What is renv?

renv (Reproducible Environment) is an R package designed to create isolated, project-specific library environments. Instead of relying on a shared system-wide R library that might change over time, renv gives each project its own separate collection of packages with specific versions.

3.2 Key Features of renv

- **Isolated project library:** renv creates a project-specific library (typically in `renv/library`) containing only the packages used by that project. This isolation ensures that updates or changes to packages in one project won't affect others.
- **Lockfile for dependencies:** When you finish installing or updating packages, `renv::snapshot()` produces a `renv.lock` file - a JSON document listing each package and its exact version and source. This lockfile is designed to be committed to version control and shared.
- **Environment restoration:** On a new machine (or when reproducing past results), `renv::restore()` installs the exact versions of packages specified in the lockfile. This creates an R package environment identical to the one that created the lockfile, provided the same R version is available. The R version is important since critical components of the R system, such as random number generation, and default factor handling policy vary between versions.

3.3 Basic renv Workflow

The typical workflow with renv involves:

```
# One-time installation of renv
install.packages("renv")

# Initialize renv for the project
renv::init() # Creates renv infrastructure

# Install project-specific packages
# ...

# Save the package state to renv.lock
renv::snapshot()
```



```
# Later or on another system...
renv::restore() # Restore packages from renv.lock
```

While renv effectively handles package dependencies, it does not address differences in R versions or system libraries. This limitation is where Docker becomes essential.

4 Docker: System-Level Reproducibility

4.1 What is Docker?

Docker is a platform that allows you to package software into standardized units called containers. A Docker container is like a lightweight virtual machine that includes everything needed to run an application: the code, runtime, system tools, libraries, and settings.

4.2 Docker's Role in Reproducibility

While renv handles R packages, Docker ensures consistency for:

- **Operating system:** The specific Linux distribution or OS version
- **R interpreter:** The exact R version
- **System libraries:** Required C/C++ libraries and other dependencies
- **Computational environment:** Memory limits, CPU configuration, etc.
- **External tools:** pandoc, LaTeX, and other utilities needed for R Markdown

By running an R Markdown project in Docker, you eliminate differences in OS or R installation as potential sources of irreproducibility. Any machine running Docker will execute the container in an identical environment.

4.3 Docker Components for R Workflows

For R-based projects, a typical Docker approach involves:

1. **Base image:** Starting from a pre-configured R image (e.g., from the Rocker project)
2. **Dependencies:** Adding system and R package dependencies
3. **Configuration:** Setting working directories and environment variables
4. **Content:** Adding project files
5. **Execution:** Defining how the project should run

The enhanced rrtools setup uses a streamlined Dockerfile based on rocker/r-ver with TinyTeX for LaTeX support:

```
FROM rocker/r-ver:4.5.0

# Prevent interactive prompts
ENV DEBIAN_FRONTEND=noninteractive

# Install minimal system dependencies
RUN apt-get update && apt-get install -y --no-install-recommends \
    pandoc \
    vim \
    git \
    curl \
    fonts-dejavu \
    && apt-get clean && rm -rf /var/lib/apt/lists/*

# Create non-root user
ARG USERNAME=analyst
RUN useradd --create-home --shell /bin/bash ${USERNAME}

# Set user R library path
ENV R_LIBS_USER=/home/${USERNAME}/R/library

# Create user R library directory and assign permissions
RUN mkdir -p /home/${USERNAME}/R/library && \
    chown -R ${USERNAME}:${USERNAME} /home/${USERNAME}/R

# Set working directory
WORKDIR /home/${USERNAME}

# Copy renv files with correct ownership
COPY --chown=${USERNAME}:${USERNAME} renv.lock ./
COPY --chown=${USERNAME}:${USERNAME} renv/activate.R ./renv/

# Switch to non-root user
USER ${USERNAME}

# Install base R packages to user library
RUN Rscript -e '.libPaths(Sys.getenv("R_LIBS_USER")); \
    install.packages(c("tinytex", "rmarkdown", "renv"), \
    repos = "https://cloud.r-project.org")'
```

```
# Install TinyTeX in user directory
RUN Rscript -e 'tinytex::install_tinytex()'

# Add TinyTeX binaries to PATH
ENV PATH=/home/${USERNAME}/.TinyTeX/bin/x86_64-linux:$PATH

# Restore R packages via renv
RUN Rscript -e '.libPaths(Sys.getenv("R_LIBS_USER")); renv::restore()'

# Default to interactive shell
CMD ["/bin/bash"]
```

Key advantages of using rocker/r-ver with TinyTeX:

- **Lightweight base:** Minimal R installation without unnecessary packages
- **TinyTeX integration:** Efficient LaTeX distribution for PDF rendering
- **Security focused:** Non-root user execution for better security
- **User-specific libraries:** Isolated package management in user directory
- **Flexible deployment:** Can be extended with additional tools as needed

Docker Compose Integration:

The setup also includes a comprehensive docker-compose.yml that provides multiple development environments:

```
# Multiple services for different workflows
services:
  r-session:    # Interactive R session
  bash:        # Bash shell access
  research:    # Automated paper rendering
  test:        # Package testing
  check:       # Package checking
```

This allows developers to choose their preferred development environment while maintaining identical package dependencies and system configuration.

5 Combining rtools, renv, and Docker: A Comprehensive Approach

5.1 Why Use All Three?

Using any single tool improves reproducibility, but combining all three provides the most comprehensive solution:

- **rrtools** provides standardized project structure and research compendium organization
- **renv** guarantees the R packages and their versions
- **Docker** guarantees the OS and R version
- **Together** they achieve end-to-end reproducibility from project organization through package dependencies to operating system consistency

This comprehensive approach creates a fully portable, well-organized research compendium that can be shared and will produce identical results across different computers while following established research best practices.

5.2 Integration Strategy with Governance Model

The recommended workflow integrates rrtools, renv, and Docker with a clear governance structure suitable for multi-developer research teams:

Project Maintainer Role (Developer 1): - Creates and maintains the research compendium structure - Manages renv environment and package dependencies
 - Updates and maintains Docker images - Reviews and approves contributor changes

Contributor Role (Other Developers): - Access the private research compendium as invited collaborators - Add analysis content, papers, and documentation using feature branches
 - Propose new package dependencies through contributions - Submit changes via pull requests from feature branches

Workflow Steps:

1. Initialize Research Compendium (Maintainer):

- Create standardized project structure with `rrtools::use_compendium()`
- Set up analysis directories with `rrtools::use_analysis()`
- Initialize renv environment with `renv::init()`
- Create Dockerfile with `rrtools::use_dockerfile()`

2. Establish Development Environment (Maintainer):

- Install required packages and develop initial analysis
- Create comprehensive tests for analytical functions
- Use `renv::snapshot()` to create initial lockfile
- Build and test Docker image locally

3. Maintain Infrastructure (Maintainer):

- Review contributor pull requests for package additions
- Update `renv.lock` by selectively incorporating new dependencies
- Rebuild Docker images when system dependencies change

- Push updated images to container registry (Docker Hub, GitHub Container Registry)

4. Collaborative Development (All Developers):

Research Compendium Files in GitHub Repository:

- **Project Structure:** DESCRIPTION, LICENSE, README.qmd (rrtools-generated)
- **Analysis Content:** Files in analysis/paper/ directory (R Markdown manuscripts)
- **Dependencies:** renv.lock (managed by maintainer), renv/activate.R
- **Infrastructure:** Dockerfile (maintained by project maintainer)
- **Code:** R/ directory (utility functions), tests/ directory
- **Documentation:** Generated README files and project documentation
- **Configuration:** .gitignore, .github/ (CI/CD workflows)

Sharing the Docker image using GitHub Container Registry:

GitHub provides GitHub Container Registry (ghcr.io) that's free for private repositories and automatically manages access permissions.

GitHub Container Registry (Recommended for Private Repos)

```
# Build the image with GitHub Container Registry URL
docker build -t ghcr.io/username/penguins_analysis:v1 .

# Login to GitHub Container Registry (using GitHub Personal Access Token)
echo $GITHUB_TOKEN | docker login ghcr.io \
  -u username --password-stdin

# Push to GitHub Container Registry (automatically private)
docker push ghcr.io/username/penguins_analysis:v1
```

Setting up GitHub Personal Access Token:

Create a Personal Access Token with the required permissions for container registry operations. The token must include write:packages and read:packages scopes, plus repo access for private repositories.

For detailed step-by-step instructions, see Appendix A: GitHub Personal Access Token Setup.

```
# Export token as environment variable (replace with your actual token)
export GITHUB_TOKEN=ghp_XXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXX

# Login and push
echo $GITHUB_TOKEN | docker login ghcr.io \
```

```
-u username --password-stdin
docker build -t ghcr.io/username/penguins_analysis:v1 .
docker push ghcr.io/username/penguins_analysis:v1
```

Note: If you get a “permission_denied” error when pushing, ensure your token includes the correct scopes (see Appendix A for details).

GitHub Container Registry Benefits:

- **Free tier:** 0.5GB storage included, no billing currently active
- **Automatic access control:** Inherits repository permissions
- **Integrated with GitHub Actions:** Seamless authentication in CI/CD
- **Simple team sharing:** Repository collaborators automatically have access
- **Package management:** Integrated with GitHub Packages ecosystem

Enhanced Docker Workflow:

The enhanced rrtools setup provides multiple approaches for working with containers, from simple Make commands to direct Docker execution. The recommended approach uses Make commands for simplicity:

```
# Build and run with Make (recommended)
make docker-build      # Build the container
make docker-r          # Interactive R session
make docker-render     # Render research paper
```

For complete Docker workflow options including Docker Compose and direct commands, see Appendix C: Docker Workflow Options.

5. Execute consistently:

- Run analyses in the Docker container for guaranteed reproducibility
- Use volume mounts to access local files while maintaining environment consistency
- Run tests within the container to verify functionality

This strategy ensures that your R Markdown documents and analyses will run identically for anyone who has access to your Docker container, regardless of their local setup.

6 Practical Example: Collaborative Research Compendium Development with Testing

The following case study demonstrates how two developers can collaborate on a research compendium using rrtools, renv, and Docker to ensure reproducibility, with integrated testing procedures to maintain code quality.

6.1 Project Scenario

Two data scientists are collaborating on an analysis of the Palmer Penguins dataset using the governance model established earlier. Developer 1 (project maintainer) will set up the initial research compendium structure using rrttools and create a basic analysis. Developer 2 (contributor) will extend the analysis with additional visualizations and propose new package dependencies. They'll use GitHub for version control and GitHub Container Registry to share the containerized environment.

Key Governance Points: - Developer 1 manages the renv environment and Docker images
- Developer 2 contributes through pull requests from their fork - Package dependency changes require Developer 1's approval and integration - Both developers use the standardized rrttools research compendium structure

6.2 Step-by-Step Implementation

6.3 Developer 1: Project Setup and Initial Analysis

Step 1: Create and Initialize the Private GitHub Repository

Developer 1 creates a new **private** GitHub repository called "penguins_analysis" and clones it locally.

Why GitHub for Private Repository + Docker Workflows (2024):

GitHub now offers excellent support for private repositories with Docker integration:

- **Free private repositories:** Unlimited private repos on free tier (since 2019)
- **GitHub Container Registry (GHCR):** Free private Docker registry with 0.5GB storage included
- **GitHub Actions:** Built-in CI/CD with excellent Docker support
- **Seamless authentication:** Registry access automatically managed with repository permissions
- **Integrated ecosystem:** Everything in one platform - code, containers, and CI/CD

Private Repository Setup:

For this collaborative scenario, we use a **private GitHub repository** to demonstrate the workflow for:

- **Team collaboration:** Controlled access where only invited team members can view and contribute
- **Proprietary research:** Working with company data, customer information, or licensed datasets

- **Sensitive research:** Medical data, personally identifiable information, or classified research
- **Commercial projects:** Business analyses, competitive intelligence, or trade secrets
- **Early development:** Preliminary research before public release or peer review
- **Institutional requirements:** When organization policies mandate private repositories

Creating a Private Repository:

1. **On GitHub.com:** Navigate to the “New repository” page
2. **Repository name:** Enter “penguins_analysis”
3. **Visibility:** Select “Private”
4. **Initialize:** Add README, .gitignore for R, and choose appropriate license
5. **Clone locally:** Use `git clone https://github.com/username/penguins_analysis.git`

Access Management for Private Repositories:

Since the repository is private, Developer 1 will need to explicitly grant access to collaborators:

- **Repository Settings → Manage access → Invite a collaborator**
- Add Developer 2’s GitHub username with “Write” access to allow branch creation and pull requests
- Developer 2 will receive an email invitation to accept

Step 2: Create Research Compendium with Enhanced rrttools Setup

Developer 1 uses the enhanced rrttools setup script to create a comprehensive research compendium:

```
# Run the enhanced rrttools setup script
./setup_rrttools.sh
```

This automated script creates a complete research compendium structure that includes:

- Enhanced R package framework with proper DESCRIPTION and NAMESPACE
- Comprehensive directory organization with data, analysis, scripts, and documentation folders
- Automated renv setup with curated package collection
- Docker integration using rocker/verse
- GitHub Actions workflows for CI/CD
- Make-based build system
- Symbolic links for easy navigation

The script automatically organizes any existing files and creates a professional research compendium structure following best practices.

Step 3: Complete the renv Environment Setup

Since the enhanced rrtools script already created the basic structure and renv configuration, Developer 1 completes the package environment setup:

```
# The setup script created setup_renv.R - run it to install packages
source("setup_renv.R")
```

This installs a comprehensive, curated collection of R packages commonly used in research projects:

- **Core tidyverse and data manipulation:** dplyr, tidyr, ggplot2, readr, etc.
- **Statistical analysis and modeling:** MASS, lme4, nlme, car, emmeans, brms, etc.
- **Data import/export:** readxl, haven, jsonlite, DBI, etc.
- **Visualization and plotting:** plotly, ggpubr, patchwork, corrplot, etc.
- **Document generation:** rmarkdown, bookdown, knitr, kableExtra, etc.
- **Package development:** devtools, usethis, testthat, roxygen2, etc.
- **Specialized analysis:** survival, psych, caret, tidymodels, etc.

The script automatically: - Installs all packages with correct dependencies - Takes a snapshot of the environment (`renv::snapshot()`) - Updates the DESCRIPTION file with dependencies

This comprehensive approach ensures most common R analysis tasks are supported without requiring additional package installations during development.

Step 4: Create Initial Analysis Paper

Developer 1 creates the analysis in the research compendium structure by editing `analysis/paper/paper.Rmd`:

```

---
title: "Palmer Penguins Analysis"
author: "Developer 1"
date: "`r Sys.Date()`"
output: pdf_document
---

```{r setup1, include=FALSE}
library(palmerpenguins)
```

# Flipper Length vs. Bill Length

```{r flipper-bill-plot1}
Developer 1's base R approach
data <- palmerpenguins::penguins
plot(data$flipper_length_mm, data$bill_length_mm,
 main = "Flipper Length vs. Bill Length",
 xlab = "Flipper Length (mm)",
 ylab = "Bill Length (mm)",
 pch = 16, col = "steelblue")
```

```

Step 5: Create Tests for Analysis Functions

While testing is uncommon in many data analysis projects, it provides significant value for reproducible research:

Why Test Data Analysis Code? - **Data integrity validation:** Ensure datasets have expected structure, ranges, and completeness - **Catch silent errors:** Detect when data changes break assumptions (e.g., missing columns, unexpected NA patterns) - **Collaboration confidence:** New team members can verify their environment setup works correctly - **Refactoring safety:** Safely improve code knowing core functionality still works - **Publication standards:** Many journals increasingly expect computational reproducibility verification - **Debugging efficiency:** Isolate whether issues stem from environment, data, or analysis logic

Types of Tests for Data Analysis: - **Data validation:** Verify data structure and content meet expectations - **Statistical sanity checks:** Ensure results fall within reasonable ranges - **Regression tests:** Confirm outputs remain consistent across environment changes - **Integration tests:** Verify the full analysis pipeline executes successfully

The Iterative Testing Process:

Testing data analysis code follows an iterative development cycle that builds confidence progressively:

1. **Start Simple:** Begin with basic data availability and structure tests that verify your dataset loads correctly and has expected dimensions. These catch fundamental setup issues early.
2. **Build Systematically:** Add tests for data types, column existence, and value ranges. Each test validates one assumption your analysis depends on.
3. **Test Incrementally:** As you develop new analysis functions, write corresponding tests before moving to the next feature. This “test-first” mindset catches issues immediately rather than during final verification.
4. **Validate Continuously:** Run tests frequently during development—after each major change, before commits, and when switching between environments. The Docker+renv setup makes this consistent across machines.
5. **Expand Coverage:** Once basic functionality works, add edge case tests, statistical validation tests, and integration tests that verify the complete analysis pipeline.

Beyond Basic Testing: The comprehensive test suite provided in the Appendix demonstrates advanced testing strategies that can be adapted for any data analysis project. These tests cover data validation, statistical relationships, visualization functions, and complete pipeline integration. Consider implementing similar comprehensive testing as your project matures, particularly for: - Long-term research projects requiring ongoing validation - Collaborative analyses where multiple team members contribute code - Production analytical pipelines that process data regularly - Academic publications where methodological rigor is essential

Developer 1 creates a test directory structure and initial tests:

```
mkdir -p tests/testthat
```

Then creates a file `tests/testthat.R`:

```
library(testthat)
library(palmerpenguins)

# Run all tests in the testthat directory
test_dir("tests/testthat")
```

And a test file `tests/testthat/test-data-integrity.R`:

```

library(testthat)
library(palmerpenguins)

test_that("penguins data is available and has expected dimensions", {
  expect_true(exists("penguins", where = "package:palmerpenguins"))
  expect_equal(ncol(palmerpenguins::penguins), 8)
  expect_gt(nrow(palmerpenguins::penguins), 300)
})

test_that("penguins data has required columns", {
  expect_true("species" %in% names(palmerpenguins::penguins))
  expect_true("bill_length_mm" %in% names(palmerpenguins::penguins))
  expect_true("flipper_length_mm" %in% names(palmerpenguins::penguins))
  expect_true("body_mass_g" %in% names(palmerpenguins::penguins))
})

```

Step 6: Create a .gitignore file

A critical aspect of reproducible projects is understanding **what should and shouldn't be tracked in version control**. Not all files created during development need to be shared—in fact, including too many files can create confusion and bloat the repository.

Files that SHOULD be tracked (committed to Git): - **Source code:** *.R, *.Rmd files containing your analysis - **Dependency specifications:** renv.lock (exact package versions), renv/activate.R (renv setup) - **Infrastructure:** Dockerfile, README.md, .gitignore - **Tests:** All files in tests/ directory that validate your analysis - **Configuration:** Any custom configuration files your analysis depends on - **Documentation:** Project documentation, methodology notes

Files that should NOT be tracked (excluded via .gitignore): - **Generated outputs:** PDFs, HTML files, plots—these are products of your code, not source materials - **Large package libraries:** renv/library/ contains downloaded packages that can be recreated from renv.lock - **Temporary files:** R session data, cache files, intermediate processing files - **Personal settings:** User-specific R configurations, local environment variables - **System artifacts:** OS-specific files, editor backup files

The principle: Track the “recipe” (code + dependencies), not the “meal” (outputs). Collaborators should run your code to generate outputs, not download pre-generated results.

Developer 1 creates a .gitignore file to exclude unnecessary files:

```

# renv - exclude downloaded packages but keep configuration
renv/library/          # Downloaded packages (recreated from renv.lock)
renv/local/            # Local package cache

```

```

renv/cellar/          # Package storage
renv/lock/            # Lock file backups
renv/python/          # Python environments
renv/staging/         # Temporary package staging

# R session files - personal and temporary
.Rhistory             # Command history (user-specific)
.RData                # Saved workspace (should start fresh)
.Ruserdata            # User session data

# Generated output files - recreated by running code
*.html                # Rendered R Markdown HTML
*.pdf                 # Rendered R Markdown PDF
*.docx                # Rendered R Markdown Word docs
output/              # Directory for analysis outputs
figures/              # Generated plots and charts
cache/                # Computation cache files

# System and editor files
.DS_Store             # macOS system files
Thumbs.db             # Windows thumbnail cache
*.tmp                 # Temporary files
*~                    # Editor backup files

```

Repository size consideration: This approach keeps the Git repository lightweight and focused. The `renv/library/` directory alone can contain hundreds of megabytes of downloaded packages, but collaborators can recreate this exactly using `renv::restore()` from the small `renv.lock` file.

Step 7: Create a Dockerfile

Developer 1 creates a Dockerfile using the enhanced rrttools setup, based on rocker/r-ver with TinyTeX for efficient LaTeX support:

```

FROM rocker/r-ver:4.5.0

# Prevent interactive prompts
ENV DEBIAN_FRONTEND=noninteractive

# Install minimal system dependencies
RUN apt-get update && apt-get install -y --no-install-recommends \
    pandoc \
    vim \

```

```

git \
curl \
fonts-dejavu \
&& apt-get clean && rm -rf /var/lib/apt/lists/*

# Create non-root user
ARG USERNAME=analyst
RUN useradd --create-home --shell /bin/bash ${USERNAME}

# Set user R library path
ENV R_LIBS_USER=/home/${USERNAME}/R/library

# Create user R library directory and assign permissions
RUN mkdir -p /home/${USERNAME}/R/library && \
    chown -R ${USERNAME}:${USERNAME} /home/${USERNAME}/R

# Set working directory
WORKDIR /home/${USERNAME}

# Copy renv files with correct ownership
COPY --chown=${USERNAME}:${USERNAME} renv.lock ./
COPY --chown=${USERNAME}:${USERNAME} renv/activate.R ./renv/

# Switch to non-root user
USER ${USERNAME}

# Install base R packages to user library
RUN Rscript -e '.libPaths(Sys.getenv("R_LIBS_USER")); \
    install.packages(c("tinytex", "rmarkdown", "renv"), \
    repos = "https://cloud.r-project.org")'

# Install TinyTeX in user directory
RUN Rscript -e 'tinytex::install_tinytex()'

# Add TinyTeX binaries to PATH
ENV PATH=/home/${USERNAME}/.TinyTeX/bin/x86_64-linux:$PATH

# Restore R packages via renv
RUN Rscript -e '.libPaths(Sys.getenv("R_LIBS_USER")); renv::restore()'

# Default to interactive shell
CMD ["/bin/bash"]

```

This enhanced Dockerfile provides:

- **Lightweight base:** Minimal R 4.5.0 installation
- **TinyTeX LaTeX:** Efficient LaTeX distribution for PDF rendering
- **Security:** Non-root user execution with proper permissions
- **User libraries:** Isolated package management in user directory
- **Interactive shell:** Default bash shell for flexible development

Step 8: Build and Push the Docker Image to GitHub Container Registry

```
# Build with GitHub Container Registry URL
docker build -t ghcr.io/username/penguins_analysis:v1 . \
  --platform=linux/amd64

# Login to GitHub Container Registry (replace with your credentials)
# Note: GITHUB_TOKEN must have 'write:packages' and 'read:packages' scopes
echo $GITHUB_TOKEN | docker login ghcr.io \
  -u username --password-stdin

# Push to GitHub Container Registry
docker push ghcr.io/username/penguins_analysis:v1
```

Important: If you encounter a “permission_denied” error during push, verify that your GitHub Personal Access Token includes the `write:packages` scope. See Appendix A for detailed instructions on creating a token with the correct permissions.

Step 9: Run tests and verify paper rendering before committing

Developer 1 runs the tests and verifies the paper renders correctly:

```
# Run all tests in the testthat directory
R -e "testthat::test_dir('tests/testthat')"

# Test that the paper renders successfully
R -e "rmarkdown::render('analysis/paper/paper.Rmd')"
```

Step 10: Commit and Push to GitHub

After confirming the tests pass and the paper renders successfully, Developer 1 commits the project files:

```
git add .
git commit -m "Initial renv setup, Docker environment, and tests"
git push origin main
```

Step 11: Communicate with Developer 2

Developer 1 provides these instructions to Developer 2:

1. Accept the invitation to join the private GitHub repository
2. Clone the repository directly as a collaborator
3. Pull the prebuilt Docker image from GitHub Container Registry
4. Run the container interactively, mounting the local repository
5. Create a new branch for feature development
6. Extend the analysis in `analysis/paper/paper.Rmd`
7. Document any new package needs (Developer 1 will manage renv updates)
8. Write tests for new functionality
9. Run tests to verify changes
10. Push changes to a feature branch on your fork and create a pull request to the main repository

Important: Developer 2 cannot directly modify `renv.lock` or update Docker images. Package dependency changes must be proposed through pull requests and will be managed by Developer 1.

6.4 Developer 2: Extending the Analysis

Step 1: Fork the Private GitHub Repository

Since the repository is **private**, Developer 2 must first be granted access by Developer 1, then create a fork for their contributions:

Fork-based Collaboration for Private Repositories

1. Developer 1 invites Developer 2 as a repository collaborator:
 - **Repository Settings** → **Manage access** → **Invite a collaborator**
 - Enter Developer 2's GitHub username (`dev2_github`)
 - Grant "Read" access (minimum required to fork private repositories)
2. Developer 2 accepts the invitation via email
3. Developer 2 forks the private repository:
 - Navigate to https://github.com/username/penguins_analysis
 - Click "Fork" button to create https://github.com/dev2_github/penguins_analysis
 - The forked repository remains private and belongs to Developer 2
4. Developer 2 clones their fork:


```
git clone \
  https://github.com/dev2_github/penguins_analysis.git
cd penguins_analysis

# Add original repository as upstream remote
git remote add upstream \
  https://github.com/username/penguins_analysis.git
```

Benefits of Fork-based Workflow for Private Repositories:

- **Clear governance:** Only Developer 1 can modify `renv.lock` and Docker images
- **Isolated development:** Developer 2's changes don't affect the main repository until approved
- **Controlled access:** Developer 2 can experiment freely in their fork
- **Professional workflow:** Follows industry best practices for controlled collaboration

Step 2: Set Up Local Development Environment

Developer 2 works with their fork and pulls the Docker image from the original repository:

```
# Already done in Step 1: clone the fork and add upstream remote
# Verify the setup
git status
git remote -v # Should show origin (fork) and upstream (original)

# Pull latest changes from upstream before starting work
git fetch upstream
git checkout main
git merge upstream/main
git push origin main # Update fork's main branch

# Login to GitHub Container Registry and pull the Docker image
echo $GITHUB_TOKEN | docker login ghcr.io -u dev2_github --password-stdin
docker pull ghcr.io/username/penguins_analysis:v1
```

Step 3: Create a Feature Branch

```
git branch body-mass-analysis
git checkout body-mass-analysis
```

Step 4: Run Docker Interactively

Developer 2 runs the container with the local repository mounted:

```
docker run --rm -it \  
  -v "$(pwd):/home/analyst/project" \  
  -v "$(pwd)/output:/home/analyst/output" \  
  -w /home/analyst/project \  
  ghcr.io/username/penguins_analysis:v1 /bin/bash
```

This approach: - Uses the renv-restored environment from the container - Mounts the local directory to /project in the container - Creates a shared output directory for generated files - Allows Developer 2 to access and modify files directly from their local machine

Step 5: Extend the Analysis

Developer 2 modifies `analysis/paper/paper.Rmd` to add a second plot for body mass vs. bill length:

```

---
title: "Palmer Penguins Analysis"
author: "Collaborative Research Team"
date: "`r Sys.Date()`"
output: pdf_document
---

```{r setup, include=FALSE}
library(ggplot2)
library(palmerpenguins)
Note: If additional packages needed (e.g., plotly, DT),
document them in pull request for Developer 1 to add
```

# Flipper Length vs. Bill Length

```{r flipper-bill-plot}
Developer 1's original base R approach
data <- palmerpenguins::penguins
plot(data$flipper_length_mm, data$bill_length_mm,
 main = "Flipper Length vs. Bill Length",
 xlab = "Flipper Length (mm)",
 ylab = "Bill Length (mm)",
 pch = 16, col = "steelblue")
```

# Body Mass vs. Bill Length

```{r mass-bill-plot}
Developer 2's contribution: Additional analysis
ggplot(palmerpenguins::penguins,
 aes(x = body_mass_g, y = bill_length_mm, color = species)) +
 geom_point() +
 theme_minimal() +
 ggtitle("Body Mass vs. Bill Length by Species")
```

```

Step 6: Create Tests for New Analysis

Developer 2 adds a new test file `tests/testthat/test-body-mass-analysis.R`:

```

test_that("body mass data is valid", {
  expect_true(all(palmerpenguins::penguins$body_mass_g > 0, na.rm = TRUE))
  expect_true(is.numeric(palmerpenguins::penguins$body_mass_g))
})

test_that("body mass correlates with bill length", {
  # Calculate correlation coefficient
  correlation <- cor(
    palmerpenguins::penguins$body_mass_g,
    palmerpenguins::penguins$bill_length_mm,
    use = "complete.obs"
  )

  # Verify correlation is a numeric value (not NA)
  expect_true(!is.na(correlation))

  # Test that the correlation is positive
  expect_true(correlation > 0)
})

```

Step 7: Run Tests and Verify Paper Rendering

Before committing, Developer 2 runs the tests and verifies the paper renders correctly with their changes:

```

# Run all tests to ensure no regressions
R -e "testthat::test_dir('tests/testthat')"

# Test that the paper renders successfully with new changes
R -e "rmarkdown::render('analysis/paper/paper.Rmd')"

```

Step 8: Commit and Push Changes to Fork

After confirming all tests pass and the paper renders successfully, Developer 2 commits and pushes the changes to their fork to prepare for a pull request:

```

git add analysis/paper/paper.Rmd \
  tests/testthat/test-body-mass-analysis.R
git commit -m \
  "Added body mass vs. bill length analysis with tests"

# Push to the feature branch on their fork
git push origin body-mass-analysis

```

Note: Developer 2 does not push directly to the main repository. All changes must go through a pull request process to maintain proper governance and code review.

Step 9: Create a Cross-Repository Pull Request

Developer 2 must submit all changes through a pull request - this is the only way to contribute to the main repository and ensures proper code review:

Developer 2 creates a pull request from their fork to the original repository:

1. Navigate to their fork: https://github.com/dev2_github/penguins_analysis
2. GitHub will typically show a banner suggesting to create a pull request after pushing a new branch
3. Click “Compare & pull request” or navigate to the Pull Requests tab and click “New pull request”
4. Ensure the pull request is configured as:
 - **Base repository:** username/penguins_analysis (the original)
 - **Base branch:** main
 - **Head repository:** dev2_github/penguins_analysis (their fork)
 - **Compare branch:** body-mass-analysis
5. Add a descriptive title: “Add body mass vs. bill length analysis”
6. In the description, include:

```
# Changes Made
- Added body mass vs. bill length scatter plot analysis
- Created comprehensive tests for body mass data validation
- Verified all existing tests continue to pass

# Testing
- All tests pass in the containerized environment
- New correlation analysis validates expected positive relationship

# Docker Environment
- Tested using `ghcr.io/username/penguins_analysis:v1` image
- No additional dependencies required

# Note
- No changes to renv.lock or Docker configuration
- Uses existing package environment
```

Benefits of Fork-based Workflow for Private Repositories: - **Controlled governance:** Only repository owner can update renv.lock and Docker images - **Isolated development:** Changes reviewed before affecting main repository - **Clear separation:** Infrastructure management vs. content contribution - **Professional workflow:** Standard practice for controlled collaboration - **Security:** Maintains strict access control over critical files

Step 10: Code Review and Merge by Developer 1

Developer 1 receives the pull request notification and conducts a thorough review:

Review Process: 1. **Examine the pull request on GitHub:** Review the code changes, commit messages, and description 2. **Test the changes locally:** Developer 1 can test the changes without affecting their main branch:

```
# Fetch the pull request branch from the fork for local testing
git fetch https://github.com/dev2_github/penguins_analysis.git \
  body-mass-analysis:review-body-mass-analysis
git checkout review-body-mass-analysis
```

3. **Verify in the Docker environment:** Test that all analyses work correctly using an interactive session:

```
# Start interactive container for comprehensive testing
docker run --rm -it \
  -v "$(pwd):/home/analyst/project" \
  -v "$(pwd)/analysis/figures:/home/analyst/output" \
  -w /home/analyst/project \
  ghcr.io/username/penguins_analysis:v1 bash

# Inside the container, run all verification steps:
# Run tests to ensure nothing breaks
R -e "testthat::test_dir('tests/testthat')"

# Test R Markdown rendering
R -e "rmarkdown::render('analysis/paper/paper.Rmd',
  output_dir='analysis/figures')"

# Verify renv environment status
R -e "renv::status()"

# Exit container when done
exit
```

4. **Manage package dependencies (if needed):** If Developer 2 has documented new package requirements in their pull request, use an interactive session for package management:

```
# Interactive session for package management
docker run --rm -it \
  -v "$(pwd):/home/analyst/project" \
  -w /home/analyst/project \
  ghcr.io/username/penguins_analysis:v1 bash

# Inside container - install packages and update lockfile
R -e "
  # Install any new packages requested by contributors
  install.packages('new_package') # Replace with actual package

  # Update the lockfile
  renv::snapshot()

  # Test that everything still works
  testthat::test_dir('tests/testthat')
"

# Exit when done
exit
```

5. **Review and merge:** If everything passes, Developer 1 merges the pull request through the GitHub interface

Post-Merge Steps: Developer 1 updates their local repository and cleans up:

```
# Switch to main branch and pull the merged changes
git checkout main
git pull origin main

# Delete the temporary review branch
git branch -d review-body-mass-analysis

# Run final verification using interactive session
docker run --rm -it \
  -v "$(pwd):/home/analyst/project" \
  -v "$(pwd)/analysis/figures:/home/analyst/output" \
  -w /home/analyst/project \
  ghcr.io/username/penguins_analysis:v1 bash
```

```
# Inside container - comprehensive final verification
R -e "
  testthat::test_dir('tests/testthat')
  rmarkdown::render('analysis/paper/paper.Rmd',
                    output_dir='analysis/figures')
  renv::status()
"

# Exit container
exit
```

Optional: Update Docker Image If the collaboration continues with more contributors, Developer 1 might consider updating the Docker image version to include any new system dependencies or optimizations:

```
# Build and push updated image (if needed)
docker build -t ghcr.io/username/penguins_analysis:v1.1 .
docker push ghcr.io/username/penguins_analysis:v1.1
```

Benefits of This Fork-based Private Repository Workflow:

- **Security:** Access is controlled and limited to invited collaborators only
- **Quality:** All changes go through review process before merging
- **Traceability:** Clear history of who made what changes and when
- **Scalability:** Multiple developers can work simultaneously on different features
- **Controlled governance:** Only Developer 1 can modify renv.lock and Docker images
- **Isolated development:** Developer 2's experiments don't affect main repository
- **Professional workflow:** Follows industry best practices for controlled collaboration
- **Clear separation:** Infrastructure management vs. content contribution responsibilities

7 Continuous Integration Extension

To further enhance the workflow, teams can set up GitHub Actions for continuous integration. This automatically runs tests in the Docker environment whenever changes are pushed, ensuring code quality and catching issues early.

Benefits of CI for Research Compendia: - **Automated testing:** Every push triggers your test suite automatically - **Environment consistency:** Tests run in the same Docker environment across all machines

- **Early error detection:** Problems are caught immediately during development - **Collaboration confidence:** Team members can see if changes break existing functionality

7.1 Key Benefits Demonstrated in This Example

This collaborative workflow demonstrates several advantages of the rrtools + renv + Docker approach:

1. **Dependency consistency:** Both developers work with identical R package versions thanks to renv.
2. **Environment consistency:** The Docker container ensures the same R version and system libraries.
3. **Code quality:** Automated tests verify that the code works as expected and catches regressions.
4. **Research compendium structure:** rrtools provides standardized organization that other researchers can easily understand.
5. **Separation of concerns:** Analysis documents remain outside the Docker image, allowing for easier collaboration.
6. **Workflow flexibility:** Developer 2 can work in the container while editing files locally.
7. **Full reproducibility:** The entire research compendium environment is captured and shareable.
8. **Continuous integration:** Automated testing ensures ongoing code quality.

For complete GitHub Actions setup instructions, workflow examples, and CI/CD configuration, see [Appendix D: GitHub Actions CI/CD Setup](#).

8 Best Practices and Considerations

8.1 When to Use This Approach

The rrtools + renv + Docker approach with testing is particularly valuable for:

- **Long-term research projects** where reproducibility over time is crucial
- **Collaborative analyses** with multiple contributors on different systems
- **Production analytical pipelines** that need to run consistently
- **Academic publications** where methods must be reproducible
- **Teaching and education** to ensure consistent student experiences
- **Complex analyses** that require rigorous testing to validate results

8.2 Tips for Efficient Implementation

1. **Keep Docker images minimal:** Include only what's necessary for reproducibility.
2. **Use specific version tags:** For both R packages and Docker base images, specify exact versions.

3. **Document system requirements:** Include notes on RAM and storage requirements.
4. **Leverage bind mounts:** Mount local directories to containers for easier development.
5. **Write meaningful tests:** Focus on validating both data integrity and analytical results.
6. **Automate testing:** Use CI/CD pipelines to automatically run tests on every change.
7. **Consider computational requirements:** Particularly for resource-intensive analyses.

8.3 Testing Strategies for R Analyses

Testing data analysis code differs from traditional software testing but provides crucial value for reproducible research:

1. **Data Validation Tests:** Ensure data has the expected structure, types, and values.
2. **Function Tests:** Verify that custom functions work as expected with known inputs and outputs.
3. **Edge Case Tests:** Check how code handles missing values, outliers, or unexpected inputs.
4. **Integration Tests:** Confirm that different parts of the analysis work correctly together.
5. **Regression Tests:** Make sure new changes don't break existing functionality.
6. **Output Validation:** Verify that final results match expected patterns or benchmarks.

While uncommon in traditional data analysis, these tests catch silent errors, validate assumptions, and provide confidence that analyses remain correct as code and data evolve.

8.4 Potential Challenges

Some challenges to be aware of:

- **Docker image size:** Images with many packages can become large
- **Learning curve:** Docker, renv, and testing frameworks require some initial learning
- **System-specific features:** Some analyses may rely on hardware features
- **Performance considerations:** Containers may have different performance characteristics
- **Test maintenance:** Tests need to be updated as the analysis evolves

9 Conclusion

Achieving full reproducibility in R requires addressing project organization, package dependencies, and system-level consistency, while ensuring code quality through testing. By combining rrttools for research compendium structure, renv for R package management, Docker for environment containerization, and automated testing for code validation, data scientists and researchers can create truly portable, reproducible, and reliable workflows.

The comprehensive approach presented in this white paper ensures that the common frustration of “it works on my machine” becomes a thing of the past. Instead, research compendia become easy to share and fully reproducible. A collaborator or reviewer can launch the Docker container and get identical results, without worrying about package versions, system setup, or project organization.

The case study demonstrates how two developers can effectively collaborate on an analysis while maintaining reproducibility and code quality throughout the project lifecycle. By integrating testing into the workflow, the team can be confident that their analysis is not only reproducible but also correct.

This strategy represents a best practice for long-term reproducibility in R, meeting the high standards required for professional data science and research documentation. The combination of standardized research compendium structure, rigorous dependency management, and containerized environments creates a robust foundation for reproducible research. By adopting this comprehensive approach, the R community can make significant strides toward the goal of fully reproducible and reliable research and analysis.

10 References

1. Marwick, B., Boettiger, C., & Mullen, L. (2018). Packaging data analytical work reproducibly using R (and friends). *The American Statistician*, 72(1), 80-88.
2. Marwick, B. (2017). rrtools: Creates a Reproducible Research Compendium. R package version 0.1.6. <https://github.com/benmarwick/rrtools>
3. Ushey, K., Wickham, H., & RStudio. (2023). renv: Project Environments. R package. <https://rstudio.github.io/renv/>
4. The Rocker Project. (2023). Docker containers for the R environment. <https://www.rocker-project.org/>
5. Wickham, H. (2023). testthat: Unit Testing for R. <https://testthat.r-lib.org/>

Appendix A: GitHub Personal Access Token Setup

This appendix provides detailed step-by-step instructions for creating a GitHub Personal Access Token with the required permissions for container registry operations.

10.1 Step-by-Step Token Creation

1. Navigate to GitHub Settings: - Go to [GitHub.com](https://github.com) and sign in - Click your profile picture (top right) → Settings - In the left sidebar: Developer settings → Personal access tokens → Tokens (classic)

- 2. Create New Token:** - Click “Generate new token” → “Generate new token (classic)”
- Add a descriptive note (e.g., “Docker Container Registry Access”) - Set expiration (recommended: 90 days for security)
- 3. Select Required Scopes** (check these boxes): - **repo** (Full control of private repositories)
- *Required for private repos* - **write:packages** (Upload packages to GitHub Package Registry)
- **read:packages** (Download packages from GitHub Package Registry) - **delete:packages** (Delete packages from GitHub Package Registry) - *Optional but recommended*
- 4. Generate and Copy Token:** - Click “Generate token” at the bottom - **Important:** Copy the token immediately - you won’t see it again - Store it securely (see security practices below)

10.2 Token Security Best Practices

- **Never commit tokens to repositories** - Use `.gitignore` to exclude files containing tokens
- **Use environment variables** - Store tokens in shell environment variables
- **Set reasonable expiration dates** - Use 30-90 day expiration for security
- **Revoke unused tokens** - Clean up tokens when no longer needed
- **Consider GitHub CLI** - Use `gh auth login` for easier management
- **Monitor token usage** - Check GitHub Settings → Developer settings → Personal access tokens for activity

10.3 Alternative: Using GitHub CLI

For easier token management, consider using GitHub CLI instead of manual tokens:

```
# Install and authenticate (handles tokens automatically)
gh auth login --scopes write:packages,read:packages,repo

# Login to container registry (automatic with gh auth)
echo $(gh auth token) | docker login ghcr.io -u $(gh api user --jq .login) --password-stdin
```

10.4 Troubleshooting Common Issues

“permission_denied: The token provided does not match expected scopes” - Verify your token includes `write:packages` and `read:packages` scopes - For private repositories, ensure `repo` scope is also selected - Create a new token with correct permissions if needed

Token not recognized: - Ensure token is properly exported: `export GITHUB_TOKEN=your_token_here`
- Verify token hasn’t expired - Check that you’re using the full token (starts with `ghp_`) 6. Horst,

A.M., Hill, A.P., & Gorman, K.B. (2020). palmerpenguins: Palmer Archipelago (Antarctica) penguin data. R package version 0.1.0. <https://allisonhorst.github.io/palmerpenguins/>
7. Marwick, B. (2016). Computational reproducibility in archaeological research: Basic principles and a case study of their implementation. *Journal of Archaeological Method and Theory*, 24(2), 424-473.

11 Appendix: Comprehensive Test Suite for Palmer Penguins Analysis

This appendix provides a complete set of tests that can be used to validate the Palmer Penguins analysis. These tests demonstrate best practices for data analysis testing and can be adapted for other projects.

11.1 Test File: tests/testthat/test-comprehensive-analysis.R

```
library(testthat)
library(palmerpenguins)
library(ggplot2)

# Test 1: Data Availability and Basic Structure
# Generic application: Verify your primary dataset loads correctly and has
# expected dimensions
# Catches: Package loading issues, file path problems, corrupted data files
test_that("Palmer Penguins dataset is available and has correct structure", {
  expect_true(exists("penguins", where = "package:palmerpenguins"))
  expect_s3_class(palmerpenguins::penguins, "data.frame")
  expect_equal(ncol(palmerpenguins::penguins), 8) # Adapt: Set expected column count
  expect_gt(nrow(palmerpenguins::penguins), 300) # Adapt: Set minimum row threshold
  expect_equal(nrow(palmerpenguins::penguins), 344) # Adapt: Set exact expected
                                                    # count if known
})

# Test 2: Required Columns Exist with Correct Types
# Generic application: Ensure your analysis depends on columns that actually
# exist with correct types
# Catches: Column name changes, type coercion issues, CSV import problems
test_that("Dataset contains required columns with expected data types", {
  df <- palmerpenguins::penguins
```

```

# Check column existence - Adapt: List columns your analysis requires
required_cols <- c("species", "island", "bill_length_mm", "bill_depth_mm",
                  "flipper_length_mm", "body_mass_g", "sex", "year")
expect_true(all(required_cols %in% names(df)))

# Check data types - Adapt: Verify types match your analysis expectations
expect_type(df$species, "integer") # Factor stored as integer
expect_type(df$bill_length_mm, "double") # Continuous measurements
expect_type(df$flipper_length_mm, "integer") # Discrete measurements
expect_type(df$body_mass_g, "integer") # Integer measurements
})

# Test 3: Categorical Variables Have Expected Levels
# Generic application: Verify factor levels for categorical variables used in
# analysis
# Catches: Missing categories, typos in factor levels, data encoding issues
test_that("Species factor has expected levels", {
  species_levels <- levels(palmerpenguins::penguins$species)
  expected_species <- c("Adelie", "Chinstrap", "Gentoo") # Adapt: Your expected
                                                         # categories
  expect_equal(sort(species_levels), sort(expected_species))
  expect_equal(length(species_levels), 3) # Adapt: Expected number of categories
  # For other datasets: Test treatment groups, regions, product types, etc.
})

# Test 4: Data Value Ranges are Domain-Reasonable
# Generic application: Verify numeric values fall within realistic ranges for
# your domain
# Catches: Data entry errors, unit conversion mistakes, outliers from
# measurement errors
test_that("Measurement values fall within reasonable biological ranges", {
  df <- palmerpenguins::penguins

  # Bill length - Adapt: Set realistic bounds for your numeric variables
  bill_lengths <- df$bill_length_mm[!is.na(df$bill_length_mm)]
  expect_true(all(bill_lengths >= 30 & bill_lengths <= 70)) # Penguin-specific
                                                            # range

  # Flipper length - Examples for other domains:
  flipper_lengths <- df$flipper_length_mm[!is.na(df$flipper_length_mm)]
  expect_true(all(flipper_lengths >= 150 & flipper_lengths <= 250))
  # Finance: stock prices > 0, percentages 0-100

```

```

# Health: age 0-120, BMI 10-80, blood pressure 50-300
# Engineering: temperatures -273+°C, pressures > 0

# Body mass
body_masses <- df$body_mass_g[!is.na(df$body_mass_g)]
expect_true(all(body_masses >= 2000 & body_masses <= 7000))
})

# Test 5: Missing Data Patterns are as Expected
# Generic application: Verify missingness patterns match your data collection
# expectations
# Catches: Unexpected data loss, systematic missingness, data pipeline failures
test_that("Missing data follows expected patterns", {
  df <- palmerpenguins::penguins

  # Total missing values should be manageable
  total_na <- sum(is.na(df))
  expect_lt(total_na, nrow(df)) # Adapt: Set acceptable threshold for missing
                                # data

  # Some variables may have expected missingness
  expect_gt(sum(is.na(df$sex)), 0) # Sex determination sometimes difficult
  # Adapt examples: Optional survey questions, historical data gaps, sensor
  # failures

  # Critical variables should be complete
  expect_equal(sum(is.na(df$species)), 0) # Primary identifier must be complete
  # Adapt: ID columns, primary keys, required fields should have no NAs
})

# Test 6: Expected Statistical Relationships Hold
# Generic application: Test known relationships between variables in your domain
# Catches: Data corruption, encoding errors, units mix-ups that break known
# patterns
test_that("Expected correlations between measurements exist", {
  df <- palmerpenguins::penguins

  # Test strong expected relationships
  correlation <- cor(df$flipper_length_mm, df$body_mass_g,
                    use = "complete.obs")
  expect_gt(correlation, 0.8) # Strong positive correlation expected
  # Adapt examples: height vs weight, price vs quality, experience vs salary

```

```

# Test weaker but expected relationships
bill_cor <- cor(df$bill_length_mm, df$bill_depth_mm, use = "complete.obs")
expect_gt(abs(bill_cor), 0.1) # Some relationship should exist
# Adapt: Education vs income, advertising vs sales, temperature vs energy use
})

# Test 7: Visualization Functions Work Correctly
# Generic application: Ensure your key plots and visualizations can be
# generated
# Catches: Missing aesthetic mappings, incompatible data types, package conflicts
test_that("Basic plots can be generated without errors", {
  df <- palmerpenguins::penguins

  # Test basic plot creation without errors
  expect_no_error({
    p1 <- ggplot(df, aes(x = flipper_length_mm, y = bill_length_mm)) +
      geom_point() +
      theme_minimal()
  })
  # Adapt: Test your key plot types - histograms, boxplots, time series, etc.

  # Test that plot object is properly created
  p1 <- ggplot(df, aes(x = flipper_length_mm, y = bill_length_mm)) +
    geom_point()
  expect_s3_class(p1, "ggplot") # Adapt: Check for your plotting framework objects
})

# Test 8: Data Filtering and Subsetting Work Correctly
# Generic application: Verify data manipulation operations produce expected results
# Catches: Logic errors in filtering, unexpected factor behaviors, indexing mistakes
test_that("Data can be properly filtered and subsetted", {
  df <- palmerpenguins::penguins

  # Test categorical filtering
  adelic_penguins <- df[df$species == "Adelie" & !is.na(df$species), ]
  expect_gt(nrow(adelic_penguins), 100) # Adapt: Expected subset size
  expect_true(all(adelic_penguins$species == "Adelie", na.rm = TRUE))
  # Adapt: Filter by treatment groups, regions, time periods, etc.

  # Test missing data handling
  complete_cases <- df[complete.cases(df), ]
  expect_lt(nrow(complete_cases), nrow(df)) # Some rows should be removed

```



```

expect_equal(sum(is.na(complete_cases)), 0) # No NAs remaining
# Adapt: Test your specific data cleaning operations
})

# Test 9: Summary Statistics are Reasonable
# Generic application: Verify computed statistics match domain knowledge
# expectations
# Catches: Calculation errors, unit mistakes, algorithm bugs, extreme outliers
test_that("Summary statistics fall within expected ranges", {
  df <- palmerpenguins::penguins

  # Test means fall within expected ranges
  mean_flipper <- mean(df$flipper_length_mm, na.rm = TRUE)
  expect_gt(mean_flipper, 190) # Adapt: Set realistic bounds for your variables
  expect_lt(mean_flipper, 210)
  # Examples: Average customer age 20-80, mean salary $30k-200k, etc.

  # Test other central tendencies
  mean_mass <- mean(df$body_mass_g, na.rm = TRUE)
  expect_gt(mean_mass, 4000)
  expect_lt(mean_mass, 5000)

  # Test variability measures are reasonable
  sd_flipper <- sd(df$flipper_length_mm, na.rm = TRUE)
  expect_gt(sd_flipper, 5) # Not zero variance
  expect_lt(sd_flipper, 30) # Not excessive variance
  # Adapt: CV should be <50%, SD should be meaningful relative to mean
})

# Test 10: Complete Analysis Pipeline Integration Test
# Generic application: Test your entire analysis workflow runs without errors
# Catches: Pipeline breaks, dependency issues, function interaction problems
test_that("Complete analysis pipeline executes successfully", {
  df <- palmerpenguins::penguins

  # Test that full workflow executes without errors
  expect_no_error({
    # Data preparation step
    clean_df <- df[complete.cases(df[c("flipper_length_mm", "bill_length_mm")]), ]

    # Statistical analysis step - Adapt: Your key analyses
    correlation_result <- cor.test(clean_df$flipper_length_mm,

```

```

                                clean_df$bill_length_mm)

# Visualization step - Adapt: Your key plots
plot_result <- ggplot(clean_df,
                      aes(x = flipper_length_mm, y = bill_length_mm)) +
  geom_point() +
  geom_smooth(method = "lm") +
  theme_minimal() +
  labs(title = "Flipper Length vs. Bill Length",
       x = "Flipper Length (mm)",
       y = "Bill Length (mm)")
})

# Adapt: Add model fitting, prediction, reporting steps as needed

# Verify analysis produces meaningful results
clean_df <- df[complete.cases(df[c("flipper_length_mm", "bill_length_mm")]), ]
correlation_result <- cor.test(clean_df$flipper_length_mm,
                              clean_df$bill_length_mm)
expect_lt(correlation_result$p.value, 0.05) # Significant result expected
# Adapt: Check model R2, prediction accuracy, convergence, etc.
})

```

11.2 Running the Tests

To run all tests in your project:

```

# Run all tests
testthat::test_dir("tests/testthat")

# Run specific test file
testthat::test_file("tests/testthat/test-comprehensive-analysis.R")

# Run tests with detailed output
testthat::test_dir("tests/testthat", reporter = "detailed")

```

11.3 Test Categories Explained

Data Validation Tests (1-5): Verify data structure, types, ranges, and missing patterns
Statistical Tests (6): Confirm expected relationships in the data **Functional Tests (7-8):**

Ensure analysis functions work correctly **Sanity Tests (9)**: Check that summary statistics are reasonable **Integration Tests (10)**: Verify the complete analysis pipeline works end-to-end

These tests provide comprehensive coverage for a data analysis project and can catch issues ranging from data corruption to environment setup problems.

Appendix B: Enhanced Directory Structure

The enhanced rrtools setup creates the following comprehensive directory structure:

```
project/
  DESCRIPTION                # Package metadata and dependencies
  LICENSE                    # Project license
  README.md                  # Project documentation
  project.Rproj              # RStudio project file
  renv.lock                  # Package dependency lockfile
  setup_renv.R               # Automated renv setup script
  Dockerfile                 # Container specification
  docker-compose.yml         # Multi-service Docker setup
  Makefile                   # Build automation (native R + Docker)
  .Rprofile                  # R startup configuration
  .dockerignore              # Docker build exclusions
  RRTOOLS_USER_GUIDE.md      # Comprehensive user documentation
  .github/workflows/         # Multiple GitHub Actions workflows
    docker-ci.yml            # Docker-based CI/CD
    r-package.yml            # R package checking
    render-paper.yml         # Automated paper rendering
  R/                          # R functions and utilities
    utils.R                  # Pre-built utility functions
  man/                       # Generated function documentation
  data/                       # Comprehensive data organization
    raw_data/                # Original, unmodified data
    derived_data/            # Processed/cleaned data
    metadata/                # Data documentation
    validation/              # Data validation scripts
    external_data/           # Third-party datasets
  analysis/                  # Research analysis
    paper/                   # Manuscript with PDF output
    figures/                 # Generated plots and charts
    tables/                  # Generated tables
    templates/               # Document templates and CSL styles
```

```

scripts/           # Working R scripts and code snippets
tests/testthat/    # Unit tests and validation
vignettes/         # Package vignettes and tutorials
inst/doc/          # Package documentation
docs/              # Additional documentation
archive/           # Archived files and old versions
[a,n,f,t,s,m,e,o,c] # Symbolic links for easy navigation

```

11.1 Key Features Explained:

Comprehensive Data Organization: - **raw_data/**: Original, unmodified datasets as received - **derived_data/**: Processed, cleaned, or transformed data - **metadata/**: Documentation about data sources, collection methods, variables - **validation/**: Scripts that verify data integrity and quality - **external_data/**: Third-party datasets or reference data

Multiple Output Formats: - **figures/**: Generated plots, charts, and visualizations - **tables/**: Generated summary tables and statistical results - **paper/**: Main manuscript and analysis documents - **templates/**: Document templates and citation style files

Professional R Package Structure: - **R/**: Custom functions and utilities - **man/**: Generated documentation for R functions - **tests/testthat/**: Unit tests and validation scripts - **vignettes/**: Long-form documentation and tutorials - **DESCRIPTION**: Package metadata and dependency specifications

Docker Orchestration: - **Dockerfile**: Main container specification - **docker-compose.yml**: Multi-service development environments - **Makefile**: Build automation supporting both native R and Docker workflows

Automated Workflows: - **.github/workflows/**: GitHub Actions for testing, checking, and rendering - **setup_renv.R**: Automated package environment setup - **RRTOOLS_USER_GUIDE.md**: Comprehensive usage documentation

Navigation Shortcuts: - **Symbolic links**: Single-letter shortcuts for easy navigation - **a** → analysis/, **n** → analysis/, **f** → figures/ - **t** → tests/, **s** → scripts/, **m** → man/ - **e** → external_data/, **o** → output/, **c** → cache/

This structure supports complex research projects while maintaining clear organization and following established research compendium principles.

Appendix C: Docker Workflow Options

The enhanced rrtools setup provides multiple approaches for working with Docker containers, each suited to different development preferences and use cases.

11.1 Option 1: Make Commands (Recommended)

The Makefile provides simplified commands that abstract complex Docker syntax:

```
# Build Docker image
make docker-build

# Interactive R session (command line users)
make docker-r

# Interactive bash session
make docker-bash

# Render research paper
make docker-render

# Run tests
make docker-test

# Package checking
make docker-check

# See all available commands
make help
```

Benefits: - **Simple syntax:** Easy-to-remember commands - **Consistent interface:** Same commands work across different projects - **Hidden complexity:** Complex Docker flags are abstracted away - **Documentation:** `make help` shows all available options

11.2 Option 2: Docker Compose Services

Docker Compose orchestrates multiple container configurations:

```
# Interactive R session
docker-compose run --rm r-session

# Bash shell access
docker-compose run --rm bash

# Automated paper rendering
docker-compose run --rm research
```

```
# Package testing
docker-compose run --rm test

# Package checking
docker-compose run --rm check
```

Docker Compose Configuration Example:

```
services:
  r-session:
    build: .
    volumes:
      - ./home/analyst/project
      - ./cache:/home/analyst/cache
    working_dir: /home/analyst/project

  bash:
    build: .
    volumes:
      - ./home/analyst/project
    working_dir: /home/analyst/project
    entrypoint: ["/bin/bash"]

  research:
    build: .
    volumes:
      - ./home/analyst/project
      - ./analysis/figures:/home/analyst/output
    working_dir: /home/analyst/project
    command: ["R", "-e", "rmarkdown::render('analysis/paper/paper.Rmd')"]
```

Benefits: - **Service orchestration:** Multiple predefined container configurations - **Volume management:** Consistent volume mounting across services - **Environment isolation:** Different services for different purposes - **Parallel execution:** Can run multiple services simultaneously

11.3 Option 3: Direct Docker Commands

For maximum control, use Docker commands directly:

```

# Basic interactive session
docker run --rm -it -v "$(pwd):/home/analyst/project" \
  ghcr.io/username/penguins_analysis:v1

# Interactive session with mounted cache
docker run --rm -it \
  -v "$(pwd):/home/analyst/project" \
  -v "$(pwd)/cache:/home/analyst/cache" \
  -w /home/analyst/project \
  ghcr.io/username/penguins_analysis:v1

# Render research paper
docker run --rm \
  -v "$(pwd):/home/analyst/project" \
  -v "$(pwd)/analysis/figures:/home/analyst/output" \
  -w /home/analyst/project \
  ghcr.io/username/penguins_analysis:v1 \
  R -e "rmarkdown::render('analysis/paper/paper.Rmd')"

# Run specific tests
docker run --rm \
  -v "$(pwd):/home/analyst/project" \
  -w /home/analyst/project \
  ghcr.io/username/penguins_analysis:v1 \
  R -e "testthat::test_file('tests/testthat/test-data-integrity.R')"

# Interactive bash session
docker run --rm -it \
  -v "$(pwd):/home/analyst/project" \
  -w /home/analyst/project \
  ghcr.io/username/penguins_analysis:v1 \
  /bin/bash

```

Common Docker Flags Explained: - **--rm**: Remove container when it exits - **-it**: Interactive terminal session - **-v**: Mount volume (host:container) - **-w**: Set working directory inside container - **--entrypoint**: Override default command

Benefits: - **Maximum flexibility**: Full control over container configuration - **Educational**: Shows exactly what's happening under the hood - **Troubleshooting**: Easier to debug when you see all options - **Portability**: Commands work on any Docker installation

11.4 Volume Mounting Strategies

Project Files:

```
# Mount entire project directory
-v "$(pwd):/home/analyst/project"
```

Output Separation:

```
# Separate outputs from source
-v "$(pwd)/analysis/figures:/home/analyst/output"
```

Cache Persistence:

```
# Persistent package cache across sessions
-v "$(pwd)/cache:/home/analyst/cache"
```

Read-only Source:

```
# Protect source files from modification
-v "$(pwd):/home/analyst/project:ro"
```

11.5 Choosing the Right Approach

Use Make Commands When: - You want simplicity and consistency - You're new to Docker - You're focusing on analysis rather than infrastructure

Use Docker Compose When: - You need multiple service configurations - You're working with a team using standardized environments - You want to define complex volume and networking setups

Use Direct Commands When: - You need maximum flexibility - You're troubleshooting container issues - You're creating custom workflows not covered by Make targets

All three approaches can be used together in the same project, depending on the specific task and user preferences.

Appendix D: GitHub Actions CI/CD Setup

GitHub Actions provides automated testing and deployment for research compendia. This appendix covers comprehensive CI/CD setup for reproducible research workflows.

11.1 Understanding GitHub Actions for Research

What is CI/CD for Research?

Continuous Integration/Continuous Deployment (CI/CD) automatically tests your research code whenever changes are made. For research compendia, this means:

- **Automated testing:** Every push triggers your test suite automatically
- **Environment consistency:** Tests run in identical Docker environments
- **Early error detection:** Problems caught immediately during development
- **Collaboration confidence:** Team members see if changes break functionality
- **Reproducibility validation:** Ensures analysis works across different systems

11.2 Step-by-Step Setup

11.2.1 Step 1: Create Workflow Directory

```
# Create the GitHub Actions directory
mkdir -p .github/workflows
```

11.2.2 Step 2: Docker-based CI Workflow

Create `.github/workflows/docker-ci.yml`:

```
name: Docker CI

on:
  push:
    branches: [ main, master ]
  pull_request:
    branches: [ main, master ]

jobs:
  build-and-test:
    runs-on: ubuntu-latest

    steps:
      - name: Checkout repository
        uses: actions/checkout@v4

      - name: Set up Docker Buildx
```

```

    uses: docker/setup-buildx-action@v3

- name: Build Docker image
  uses: docker/build-push-action@v5
  with:
    context: .
    push: false
    tags: ${GITHUB_REPOSITORY}:latest
    cache-from: type=gha
    cache-to: type=gha,mode=max

- name: Run tests in container
  run: |
    docker run --rm -v $PWD:/home/analyst/project \
      ${GITHUB_REPOSITORY}:latest \
      R -e "testthat::test_dir('tests/testthat')"

- name: Render research paper
  run: |
    docker run --rm -v $PWD:/home/analyst/project \
      -v $PWD/analysis/figures:/home/analyst/output \
      ${GITHUB_REPOSITORY}:latest \
      R -e "rmarkdown::render('analysis/paper/paper.Rmd')"

- name: Upload rendered paper
  uses: actions/upload-artifact@v4
  if: success()
  with:
    name: research-paper
    path: analysis/paper/paper.pdf

```

11.2.3 Step 3: R Package Check Workflow

Create `.github/workflows/r-package.yml`:

```

name: R Package Check

on:
  push:
    branches: [ main, master ]
  pull_request:

```

```

    branches: [ main, master ]

jobs:
  R-CMD-check:
    runs-on: ${ matrix.config.os }

    name: ${ matrix.config.os } (${ matrix.config.r })

    strategy:
      fail-fast: false
      matrix:
        config:
          - {os: ubuntu-latest,    r: 'release'}
          - {os: macOS-latest,    r: 'release'}
          - {os: windows-latest,  r: 'release'}

    env:
      GITHUB_PAT: ${ secrets.GITHUB_TOKEN }
      R_KEEP_PKG_SOURCE: yes

    steps:
      - uses: actions/checkout@v4

      - uses: r-lib/actions/setup-pandoc@v2

      - uses: r-lib/actions/setup-r@v2
        with:
          r-version: ${ matrix.config.r }
          http-user-agent: ${ matrix.config.http-user-agent }
          use-public-rspm: true

      - uses: r-lib/actions/setup-renv@v2

      - name: Install system dependencies
        if: runner.os == 'Linux'
        run: |
          sudo apt-get update
          sudo apt-get install -y \
            libcurl4-openssl-dev \
            libssl-dev \
            libxml2-dev

```

```
- uses: r-lib/actions/check-r-package@v2
  with:
    upload-snapshots: true
```

11.2.4 Step 4: Automated Paper Rendering

Create `.github/workflows/render-paper.yml`:

```
name: Render Research Paper

on:
  workflow_dispatch: # Manual trigger
  push:
    branches: [ main, master ]
    paths:
      - 'analysis/paper/**'
      - 'analysis/data/**'
      - 'R/**'
      - 'data/**'

jobs:
  render:
    runs-on: ubuntu-latest

    steps:
      - name: Checkout repository
        uses: actions/checkout@v4

      - name: Set up Docker Buildx
        uses: docker/setup-buildx-action@v3

      - name: Build Docker image
        uses: docker/build-push-action@v5
        with:
          context: .
          push: false
          tags: paper-render:latest
          cache-from: type=gha
          cache-to: type=gha,mode=max

      - name: Render paper in container
```

```

run: |
  docker run --rm \
    -v $PWD:/home/analyst/project \
    -v $PWD/analysis/figures:/home/analyst/output \
    paper-render:latest \
    R -e "rmarkdown::render('analysis/paper/paper.Rmd')"

- name: Upload rendered paper
  uses: actions/upload-artifact@v4
  with:
    name: research-paper-${{ github.sha }}
    path: |
      analysis/paper/paper.pdf
      analysis/figures/*.png
      analysis/figures/*.jpg
    retention-days: 30

```

11.2.5 Step 5: Container Registry Integration

Create `.github/workflows/container-publish.yml`:

```

name: Build and Push Container

on:
  push:
    branches: [ main ]
    tags: [ 'v*' ]
  pull_request:
    branches: [ main ]

env:
  REGISTRY: ghcr.io
  IMAGE_NAME: ${ github.repository }

jobs:
  build-and-push:
    runs-on: ubuntu-latest
    permissions:
      contents: read
      packages: write

```

```

steps:
  - name: Checkout repository
    uses: actions/checkout@v4

  - name: Set up Docker Buildx
    uses: docker/setup-buildx-action@v3

  - name: Log in to Container Registry
    if: github.event_name != 'pull_request'
    uses: docker/login-action@v3
    with:
      registry: ${ env.REGISTRY }
      username: ${ github.actor }
      password: ${ secrets.GITHUB_TOKEN }

  - name: Extract metadata
    id: meta
    uses: docker/metadata-action@v5
    with:
      images: ${ env.REGISTRY }/${ env.IMAGE_NAME }
      tags: |
        type=ref,event=branch
        type=ref,event=pr
        type=semver,pattern={{version}}
        type=semver,pattern={{major}}.{{minor}}

  - name: Build and push Docker image
    uses: docker/build-push-action@v5
    with:
      context: .
      platforms: linux/amd64,linux/arm64
      push: ${ github.event_name != 'pull_request' }
      tags: ${ steps.meta.outputs.tags }
      labels: ${ steps.meta.outputs.labels }
      cache-from: type=gha
      cache-to: type=gha,mode=max

```

11.3 Workflow Explanations

11.3.1 Docker CI Workflow Features:

- **Build Testing:** Ensures Docker image builds with latest changes
- **Comprehensive Testing:** Runs R package tests and renders paper in container
- **Artifact Generation:** Saves rendered papers as downloadable artifacts
- **Caching:** Uses GitHub Actions cache for faster builds

11.3.2 R Package Check Features:

- **Multi-platform Testing:** Tests on Ubuntu, macOS, and Windows
- **R CMD Check:** Comprehensive package validation
- **renv Integration:** Automatically restores package environment
- **System Dependencies:** Installs required system libraries

11.3.3 Paper Rendering Features:

- **Selective Triggering:** Only runs when relevant files change
- **Manual Execution:** Can be triggered manually via GitHub interface
- **Artifact Storage:** Saves PDFs and figures with retention policy
- **Path-based Triggers:** Responds to changes in analysis files

11.3.4 Container Publishing Features:

- **Automated Building:** Builds on pushes and tags
- **Multi-architecture:** Supports AMD64 and ARM64 platforms
- **Semantic Versioning:** Automatic tagging based on git tags
- **Security:** Uses built-in GitHub token for authentication

11.4 Authentication and Permissions

11.4.1 Built-in GITHUB_TOKEN:

The built-in GITHUB_TOKEN automatically provides: - Read access to repository contents - Write access to GitHub Packages (when permissions are set) - No manual setup required

11.4.2 Setting Repository Permissions:

1. **Repository Settings** → **Actions** → **General**
2. **Workflow permissions:** Choose “Read and write permissions”
3. **Allow GitHub Actions to create and approve pull requests:** Enable if needed

11.4.3 Using Personal Access Tokens (Advanced):

For broader permissions, create repository secrets:

1. **Repository Settings** → **Secrets and variables** → **Actions**
2. **New repository secret:** Add GHCR_TOKEN with Personal Access Token
3. **Reference in workflow:** password: `${{ secrets.GHCR_TOKEN }}`

11.5 Integration with Collaborative Workflow

11.5.1 Pull Request Integration:

When Developer 2 submits a pull request: 1. GitHub automatically triggers CI workflows 2. Tests run in clean environment identical to production 3. Results displayed directly in pull request interface 4. Merge can be blocked if tests fail

11.5.2 Branch Protection Rules:

Enable in **Repository Settings** → **Branches**: - **Require status checks:** Force CI to pass before merging - **Require branches to be up to date:** Ensure latest code is tested - **Include administrators:** Apply rules to all users

11.6 Monitoring and Troubleshooting

11.6.1 Viewing Workflow Results:

1. **Repository** → **Actions** tab
2. Click specific workflow run to see details
3. Expand steps to see detailed logs
4. Download artifacts (rendered papers, test results)

11.6.2 Common Issues and Solutions:

Docker Build Failures: - Check Dockerfile syntax - Verify all COPY paths exist - Ensure base image is accessible

renv Restore Failures: - Verify renv.lock is committed - Check for platform-specific packages - Consider using RSPM for faster installs

Permission Errors: - Verify GITHUB_TOKEN permissions - Check repository secrets configuration - Ensure workflows have necessary permissions

11.6.3 Performance Optimization:

Caching Strategies: - Docker layer caching with `cache-from/cache-to` - renv package caching with `r-lib/actions/setup-renv` - Artifact caching for large datasets

Parallel Execution: - Run tests and documentation in parallel jobs - Use matrix strategies for multi-platform testing - Conditional execution based on changed files

This comprehensive CI/CD setup ensures that research compendia remain reproducible, tested, and deployment-ready throughout the development lifecycle.